

This project gathered data from WeRateDogs account on Twitter. There were three different datasets that needed to be assessed for quality and tidiness. These datasets were tweet-json.txt, image-predictions.tsv and twitter-archive-enhanced.csv.

## Quality Issue 1, 8 and 11

There are three occasions where we would need to convert a column data type from integer (int) to strings (str). To reduce the time needed to code for every single column as well as repeat long code repetition, we created a function to deal with this quality issue. The code is as shown below:

```
# We will write a def function here to make all other conversions simpler by calling it
def conv_to_str(dataset,column):
    dataset[column] = dataset[column].astype(str)

conv_to_str(twitter_archive_clean,'tweet_id')
```

We test our cleaning attempt with the .type() method like so:

### Test

```
type(twitter_archive_clean.tweet_id[0])
str
```

## Quality Issue 2

We also convert the timestamp in the twitter archive dataframe from a string to a dateTime type. This is to make sure that dateTime analysis can be computed successfully. The code to enable this is executed using pd.to\_datetime() and also tested with the .type() method

### Code

```
In [51]: twitter_archive_clean['timestamp'] = pd.to_datetime(twitter_archive_clean['timestamp'])
```

### Test

```
In [52]: type(twitter_archive_clean['timestamp'][0])
Out[52]: pandas._libs.tslibs.timestamps.Timestamp
```

## Quality Issue 3 and 4

With this project, we required only original tweets and not retweets so we dropped all columns that had anything to do with retweets. These columns were found in the twitter\_archive dataset during assessment. The respective columns were dropped using the .drop() method. One way of confirming if our method worked is to check the column information the twitter archive dataset holds. This is also done by using the .info() method. Both methods are executed as shown;

## Code

```
In [53]: twitter_archive_clean = twitter_archive_clean.drop(columns=
        ['retweeted_status_id',
         'retweeted_status_user_id',
         'retweeted_status_timestamp']
        )
```

## Test

```
In [54]: twitter_archive_clean.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   tweet_id              2356 non-null   object
1   in_reply_to_status_id  78 non-null     float64
2   in_reply_to_user_id    78 non-null     float64
3   timestamp             2356 non-null   datetime64[ns, UTC]
4   source                2356 non-null   object
5   text                 2356 non-null   object
6   expanded_urls         2297 non-null   object
7   rating_numerator       2356 non-null   int64
8   rating_denominator     2356 non-null   int64
9   name                 2356 non-null   object
10  doggo                 2356 non-null   object
11  floofer               2356 non-null   object
12  pupper               2356 non-null   object
13  puppo                 2356 non-null   object
dtypes: datetime64[ns, UTC](1), float64(2), int64(2), object(9)
memory usage: 257.8+ KB
```

We repeated the same action for other column features we won't need in our analysis

## Code

```
In [55]: twitter_archive_clean = twitter_archive_clean.drop(columns=
        ['in_reply_to_status_id', 'in_reply_to_user_id',
         'expanded_urls',
         'source' ]
        )
```

## Test

```
In [56]: twitter_archive_clean.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   tweet_id              2356 non-null   object
1   timestamp             2356 non-null   datetime64[ns, UTC]
2   text                 2356 non-null   object
3   rating_numerator       2356 non-null   int64
4   rating_denominator     2356 non-null   int64
5   name                 2356 non-null   object
6   doggo                 2356 non-null   object
7   floofer               2356 non-null   object
8   pupper               2356 non-null   object
9   puppo                 2356 non-null   object
dtypes: datetime64[ns, UTC](1), int64(2), object(7)
memory usage: 184.2+ KB
```

## Quality Issue 5 and 7

Column names should be made easy to understand at first glance by the analyst so there isn't the need for constant referencing. Constant referencing wastes time. A few columns in the image prediction dataset would need their column names renamed. For this purpose we would deploy the `.rename()` method and check with the `info()` method as shown;

**Code**

```
In [57]: image_predictions_clean.rename(columns={'p1':'algor_1',
                                             'p2':'algor_2',
                                             'p3':'algor_3',
                                             'p1_conf':'algor_1_conf',
                                             'p2_conf':'algor_2_conf',
                                             'p3_conf':'algor_3_conf',
                                             'p1_dog':'algor_1_bool',
                                             'p2_dog':'algor_2_bool',
                                             'p3_dog':'algor_3_bool'}, inplace=True);
```

**Test**

```
In [58]: image_predictions_clean.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   tweet_id        2075 non-null   int64
1   jpg_url         2075 non-null   object
2   img_num         2075 non-null   int64
3   algor_1         2075 non-null   object
4   algor_1_conf    2075 non-null   float64
5   algor_1_bool    2075 non-null   bool
6   algor_2         2075 non-null   object
7   algor_2_conf    2075 non-null   float64
8   algor_2_bool    2075 non-null   bool
9   algor_3         2075 non-null   object
10  algor_3_conf     2075 non-null   float64
11  algor_3_bool     2075 non-null   bool
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 152.1+ KB
```

We also name the 'id' column in the tweet json dataset to 'tweet\_id' so we could merge by that column with other dataframes

**Code**

```
In [62]: tweet_df_clean.rename(columns={'id':'tweet_id'}, inplace=True)
```

**Test**

```
In [63]: tweet_df_clean.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2354 entries, 0 to 2353
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   tweet_id        2354 non-null   int64
1   retweet_count    2354 non-null   int64
2   favorite_count   2354 non-null   int64
dtypes: int64(3)
memory usage: 55.3 KB
```

## Quality Issue 6

It is important that all null values are correctly represented otherwise they could be categorized as something which might affect our the quality of our analysis. It can be seen from the twitter archive dataset that the null values are represented by 'None' instead of 'NaN'. We correct this issue by replacing all 'None' values with 'NaN' using numpy's np.nan and check for nulls using the .isnull() method.

**Code**

```
In [59]: twitter_archive_clean.name = twitter_archive_clean.name.replace('None', value = np.nan)
```

**Test**

```
In [60]: twitter_archive_clean.name.isnull().sum()
```

```
Out[60]: 745
```

## Quality Issue 9

Some breed names in the predictions column in the image prediction dataset starts with an upper-case while others start with a lower case showing non-uniformity. These columns are represented by algor\_1, algor\_2 and algor\_3. All names starting with lower-case letters were converted to start with upper-case by deploying the .str.title() method. This was tested by calling a sample and visually assessing.

### Code

```
In [67]: image_predictions_clean.algor_1 = image_predictions_clean.algor_1.str.title()
image_predictions_clean.algor_2 = image_predictions_clean.algor_2.str.title()
image_predictions_clean.algor_3 = image_predictions_clean.algor_3.str.title()
```

### Test

```
In [68]: image_predictions_clean.sample(5)
```

Out[68]:

jpg_url	img_num	algor_1	algor_1_conf	algor_1_bool	algor_2	algor_2_conf	algor_2_bool	algor_3	algor_3_conf	algor_3_bool
gWIAAwOUy.jpg	1	Shih-Tzu	0.225848	True	Norfolk_Terrier	0.186873	True	Irish_Terrier	0.106987	True
LWEAEEnSk7.jpg	1	Microwave	0.981946	False	Rotisserie	0.007472	False	Television	0.005881	False
7UcAANzmK.jpg	1	Soft-Coated_Wheaten_Terrier	0.403496	True	Cocker_Spaniel	0.135164	True	Golden_Retriever	0.088719	True
odXAAAJ1de.jpg	1	Hamster	0.132440	False	Toy_Poodle	0.123962	True	Bubble	0.056212	False
WMMAAoG5.jpg	1	Llama	0.379624	False	Triceratops	0.162761	False	Hog	0.084251	False

## Quality Issue 10

According to WeRateDogs, all dogs are awesome so we cant have a rating lower than the perfect score, 10/10. This means we need to make sure at least all ratings numerator in the twitter archive dataset is at least 10. We first query tha dataset for all ratings numerator values less than 10 and add 10 to them using the .add() method. By deploying the .value\_counts() method, we can see that our lowest value is now 10.

### Code

```
In [70]: twitter_archive_clean.rating_numerator[twitter_archive_clean['rating_numerator'] <= 10 ] =
twitter_archive_clean.rating_numerator[twitter_archive_clean['rating_numerator'] <= 10 ].add(10)

C:\Users\dell\AppData\Local\Temp\ipykernel_6440\3417593879.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
twitter_archive_clean.rating_numerator[twitter_archive_clean['rating_numerator'] <= 10 ] = twitter_archive_clean.rating_numerator[twitter_archive_clean['rating_numerator'] <= 10 ].add(10)
```

```
In [289]: twitter_archive_clean.rating_numerator.value_counts()
```

Out[289]:

```
12    567
11    473
20    462
13    370
19    158
18    102
14     71
17     56
15     39
16     32
420     2
10     2
75     2
45     1
60     1
44     1
121    1
143    1
99     1
```



## Tidiness Issue 1

To be able to do relational analysis between each dataframe easily we would need to merge the three dataframes. The `.merge()` method was called here. We first merged the image predictions and twitter archive into one master dataframe by the 'tweet\_id' in both dataframes. Then we merged that master dataframe with the tweet json dataframe by the same 'tweet\_id'.

### Code

```
In [73]: master_data_clean = []

master_data_clean = pd.merge( image_predictions_clean, twitter_archive_clean, on='tweet_id', how="left")

In [74]: master_data_clean = pd.merge( master_data_clean, tweet_df_clean, on='tweet_id', how="left")
```

### Test

```
In [75]: master_data_clean.shape

Out[75]: (2075, 23)

In [76]: master_data_clean.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2075 entries, 0 to 2074
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   tweet_id              2075 non-null   object
1   jpg_url               2075 non-null   object
2   img_num               2075 non-null   int64
3   algor_1              2075 non-null   object
4   algor_1_conf         2075 non-null   float64
5   algor_1_bool         2075 non-null   bool
6   algor_2              2075 non-null   object
7   algor_2_conf         2075 non-null   float64
8   algor_2_bool         2075 non-null   bool
9   algor_3              2075 non-null   object
10  algor_3_conf         2075 non-null   float64
11  algor_3_bool         2075 non-null   bool
12  timestamp             2075 non-null   datetime64[ns, UTC]
13  text                  2075 non-null   object
14  rating_numerator      2075 non-null   int64
15  rating_denominator    2075 non-null   int64
16  name                  1497 non-null   object
17  doggo                 2075 non-null   object
18  floofer               2075 non-null   object
19  pupper                2075 non-null   object
20  puppo                 2075 non-null   object
21  retweet_count         2073 non-null   float64
22  favorite_count        2073 non-null   float64
dtypes: bool(3), datetime64[ns, UTC](1), float64(5), int64(3), object(11)
memory usage: 346.5+ KB
```

## Tidiness Issue 2

The doggo, floofer, pupper and puppo columns in twitter\_archive should be categories in one column. We use the `pd.melt()` method to create a new column called 'stage' in the tweet json dataframe to house the dog stages and melt the four individual stages into a new 'stage' column in the master data and finally delete existing four stages column.

### Code

```
In [77]: df = pd.melt(master_data_clean, id_vars=['tweet_id'], value_vars=['doggo', 'floofer', 'pupper', 'puppo'], value_name='stage')

master_data_clean['stage'] = df['stage']
```

**Test**

```
[386]: master_data_clean.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 33200 entries, 0 to 33199
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   tweet_id                             33200 non-null  object
1   jpg_url                              33200 non-null  object
2   img_num                              33200 non-null  int64
3   algor_1                             33200 non-null  object
4   algor_1_conf                         33200 non-null  float64
5   algor_1_bool                        33200 non-null  bool
6   algor_2                             33200 non-null  object
7   algor_2_conf                        33200 non-null  float64
8   algor_2_bool                        33200 non-null  bool
9   algor_3                             33200 non-null  object
10  algor_3_conf                        33200 non-null  float64
11  algor_3_bool                        33200 non-null  bool
12  in_reply_to_status_id               368 non-null    float64
13  in_reply_to_user_id                368 non-null    float64
14  timestamp                          33200 non-null  object
15  source                             33200 non-null  object
16  text                               33200 non-null  object
17  retweeted_status_id                1296 non-null    float64
18  retweeted_status_user_id           1296 non-null    float64
19  retweeted_status_timestamp          1296 non-null    object
20  expanded_urls                      33200 non-null  object
21  rating_numerator                   33200 non-null  int64
22  rating_denominator                 33200 non-null  int64
23  name                               23952 non-null  object
24  retweet_count                      33168 non-null  float64
25  favorite_count                     33168 non-null  float64
26  stage                              1336 non-null    object
dtypes: bool(3), float64(9), int64(3), object(12)
memory usage: 6.4+ MB
```

## Tidiness Issue 3

We would finally extract day, month and year from the timestamp into individual columns. We use the .dt method to create new variables called 'day', 'month' and 'year' and populate it with new columns in the master dataframe with the same names.

**Code**

```
In [102]: day = master_data_clean['timestamp'].dt.day
month = master_data_clean['timestamp'].dt.month
year = master_data_clean['timestamp'].dt.year
```

```
In [104]: master_data_clean['day'] = day
master_data_clean['month'] = month
master_data_clean['year'] = year
```

**Test**

```
In [103]: master_data_clean.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2075 entries, 0 to 2074
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   tweet_id              2075 non-null   object
1   jpg_url               2075 non-null   object
2   img_num               2075 non-null   int64
3   algor_1               2075 non-null   object
4   algor_1_conf          2075 non-null   float64
5   algor_1_bool          2075 non-null   bool
6   algor_2               2075 non-null   object
7   algor_2_conf          2075 non-null   float64
8   algor_2_bool          2075 non-null   bool
9   algor_3               2075 non-null   object
10  algor_3_conf          2075 non-null   float64
11  algor_3_bool          2075 non-null   bool
12  timestamp              2075 non-null   datetime64[ns, UTC]
13  text                   2075 non-null   object
14  rating_numerator       2075 non-null   int64
15  rating_denominator     2075 non-null   int64
16  name                   1497 non-null   object
17  doggo                  2075 non-null   object
18  floofer                2075 non-null   object
19  pupper                 2075 non-null   object
20  puppo                  2075 non-null   object
21  retweet_count          2073 non-null   float64
22  favorite_count         2073 non-null   float64
23  stage                  2075 non-null   object
dtypes: bool(3), datetime64[ns, UTC](1), float64(5), int64(3), object(12)
memory usage: 362.7+ KB
```